

CS-6240 HW-1 Report

Ajay Baban Kauthale(ajayk@ccs.neu.edu)

Weather Data Source Code

Source code is in "HW01/1. Source Code" directory

Weather Data Results

A) Running times without Fibonacci

| Model | Min | Max | Avg |
|-------------|------|------|------|
| Sequential | 2624 | 6224 | 3088 |
| No Lock | 1674 | 5264 | 2083 |
| Coarse Lock | 2103 | 5776 | 2512 |
| Fine Lock | 1667 | 6437 | 2271 |
| No Sharing | 1819 | 5176 | 2364 |

B) Running times with Fibonacci

| Model | Min | Max | Avg |
|-------------|-------|-------|-------|
| Sequential | 12730 | 16553 | 13262 |
| No Lock | 6600 | 11598 | 7394 |
| Coarse Lock | 8135 | 12000 | 8603 |
| Fine Lock | 6585 | 11598 | 7535 |
| No Sharing | 6985 | 11905 | 8349 |

C) Number of threads: 4 (dual core processor, 4 logical CPU's)

Speed Up:

| Model | Speed up without delay | Speed up with delay |
|-------------|------------------------|---------------------|
| No Lock | 3.33 | 2.46 |
| Coarse Lock | 3.13 | 2.34 |
| Fine Lock | 3.56 | 2.41 |
| No Sharing | 2.95 | 2.26 |

D) Questions:

1. **Which** program version (SEQ, NO-LOCK, COARSE-LOCK, FINE-LOCK, NO-SHARING) would you normally expect to finish fastest and **why**? Do the experiments confirm your expectation? If not, try to **explain** the reasons.

- ⇒ **No Lock** should be the fastest one.
- ⇒ The reason I think No Lock should be fastest is because it does not lock data in any way. Because of no data locking each thread can access data simultaneously. Although it does not guarantee correct output, it is fastest one which is confirmed by above results and it meets the expectation.

2. **Which** program version (SEQ, NO-LOCK, COARSE-LOCK, FINE-LOCK, NO-SHARING) would you normally expect to finish slowest and **why**? Do the experiments confirm your expectation? If not, try to explain the reasons.

- ⇒ **Sequential** should be the slowest one.
- ⇒ The reason I think Sequential should be slowest one is because it is not concurrent and single process reading and processing data sequentially. Hence, it should be slowest as compared to any other multithreaded version which is confirmed by my experiment.

3. Compare the temperature averages returned by each program version. Report if any of them is **incorrect** or if any of the programs **crashed** because of concurrent accesses.

- ⇒ During the experiment, no program is crashed. Only discrepancy I noticed is whenever I start other applications (browser etc.) while running my experiment all running times go high. Especially **No Lock** gives different result for every other run.
- ⇒ This seems expected since applications might be competing for the system resources. Also, **No Lock** do not have synchronization in place which is causing different results.

4. Compare the running times of SEQ and COARSE-LOCK. Try to explain **why** one is slower than the other. (Make sure to consider the results of both B and C—this might support or refute a possible hypothesis.)

- ⇒ **Sequential** should be slower than **Coarse Lock** because as we discussed above Sequential process all data one by one in single process which is time consuming. On the other hand, Coarse Lock uses concurrent threads to process that data making it faster than that of the Sequential.
- ⇒ Experiment confirms my understanding and hypothesis is true. The results are as follows

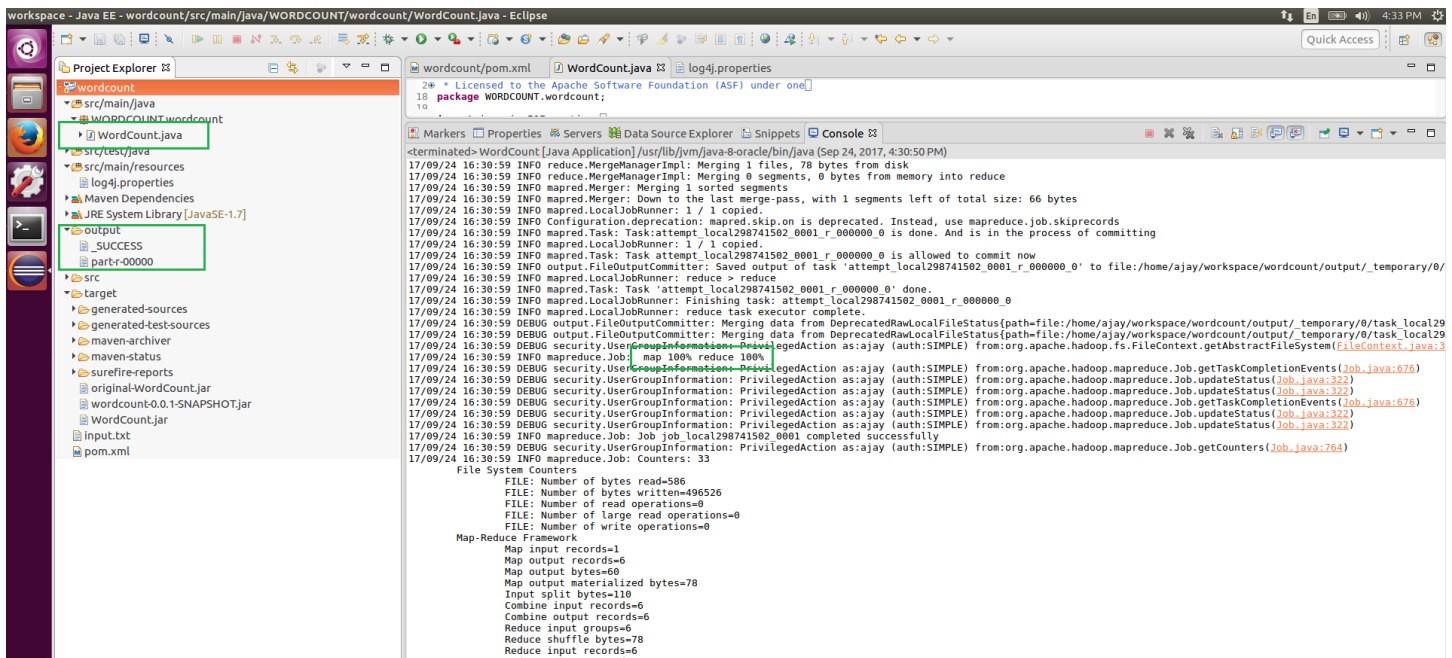
| Model | Avg time without Fibonacci | Avg time with Fibonacci |
|-------------|----------------------------|-------------------------|
| Sequential | 3088 | 13262 |
| Coarse Lock | 2512 | 8603 |

5. **How** does the higher computation cost in part C (additional Fibonacci computation) affect the difference between COARSE-LOCK and FINE-LOCK? Try to **explain** the reason.

- ⇒ Both **Coarse Lock** and **Fine Lock** running times increased with Fibonacci computation.
- ⇒ For the **Coarse Lock**, the difference is slightly greater as compared to **Fine Lock** since all data is locked in **Coarse Lock** instead of station specific data.
- ⇒ For the **Fine Lock**, the station data is locked instead of all data. Therefore, insertion and processing is faster as compared to the **Coarse Lock**.

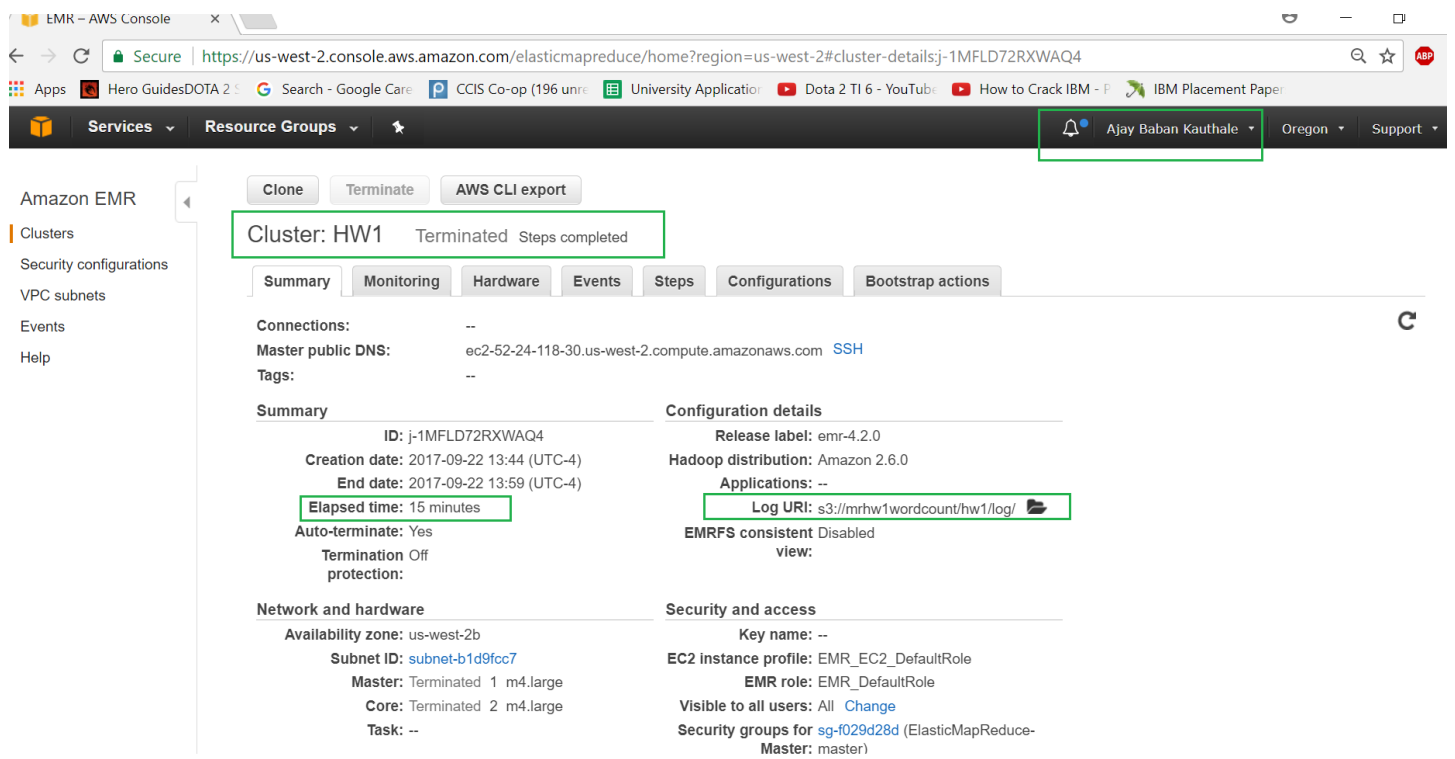
Word Count Local Execution (UBUNTU)

Following screenshot contains both directory structure and console output (> 20 lines)



Word Count AWS Execution

Following screenshot contains cluster screenshot with three machines (1 master and 2 workers) and amazon S3 storage.



Amazon S3

S3 Management Console

Secure | https://s3.console.aws.amazon.com/s3/buckets/mrhw1wordcount/hw1/?region=us-west-2&tab=overview

AppsHero GuidesDOTA 2Search - Google CareCCIS Co-op (196 unreUniversity ApplicationDota 2 TI 6 - YouTubeHow to Crack IBM - PIBM Placement Paper

ServicesResource GroupsAjay Baban KauthaleGlobalSupport

Amazon S3 > mrhw1wordcount / hw1

Overview

Q Type a prefix and press Enter to search. Press ESC to clear.

Upload

Create folder

More

Versions

Hide

Show

US West (Oregon)

| Viewing 1 to 5 | | | | |
|--------------------------|---------------|--------------------------|---------|---------------|
| <input type="checkbox"/> | Name | Last modified | Size | Storage class |
| <input type="checkbox"/> | input | -- | -- | -- |
| <input type="checkbox"/> | log | -- | -- | -- |
| <input type="checkbox"/> | output | -- | -- | -- |
| <input type="checkbox"/> | output1 | -- | -- | -- |
| <input type="checkbox"/> | WordCount.jar | Sep 22, 2017 12:25:52 PM | 39.0 MB | Standard |